

importing modules

```
import pandas as pd
import requests as req
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.pipeline import Pipeline
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import precision_score, recall_score, f1_score
from sklearn.model_selection import ParameterGrid
from sklearn.preprocessing import LabelEncoder
import re
```

*****PROBLEM 1 – Reading the data*****

urls to the datasets

```
positive_baits="https://raw.githubusercontent.com/pfrcks/clickbait-  
detection/master/clickbait"  
negative_baits="https://raw.githubusercontent.com/pfrcks/clickbait-  
detection/master/not-clickbait"
```

A function to fetch the data from a URL

```
def my_read_function(url,label):  
    response=req.get(url)  
    text = response.text  
    lines = text.split('\n')  
    df=pd.DataFrame({'baits': lines})  
    df.index.name = 'Index'  
    df['Label'] = label  
    #df.to_csv('baits_data.csv', index=False)  
    return df
```

Read the positive and negative datasets

```
positive_dataset = my_read_function(positive_baits,'clickbait')
negative_dataset = my_read_function(negative_baits,'not-clickbait')
```

```
#reading from the positive_dataset/click_baits_dataset
df=pd.DataFrame(positive_dataset)
df
```

	baits	Label
Index		
0	Man repairs fence to contain dog, hilarity ens...	clickbait
1	Long-Term Marijuana Use Has One Crazy Side Eff...	clickbait
2	The water from his ear trickles into the bucke...	clickbait
3	You'll Never Guess What Nick Jonas Does in the...	clickbait
4	How Cruise Liners Fill All Their Unsold Cruise...	clickbait
...
810	OITNB's Taylor Schilling and Carrie Brownstein...	clickbait
811	Researchers have discovered the average penis ...	clickbait
812	Why it may be smart to wait to put on sunscree...	clickbait
813	What state has highest rate of rape in the cou...	clickbait
814		clickbait

[815 rows x 2 columns]

```
#reading from the positive_dataset/click_baits_dataset
df=pd.DataFrame(negative_dataset)
df
```

	baits	
Label		
Index		
0	Congress Slips CISA Into a Budget Bill That's ...	not-clickbait
1	DUI Arrest Sparks Controversy	not-clickbait
2	It's unconstitutional to ban the homeless from...	not-clickbait
3	A Government Error Just Revealed Snowden Was t...	not-clickbait
4	A toddler got meningitis. His anti-vac parents...	not-clickbait
...
.		
1570	Loophole means ecstasy and loads of other drug...	not-clickbait

```

1571      Astronomers Watch a Supernova and See Reruns  not-
clickbait
1572  In Indian Rapists' Neighborhood, Smoldering An...  not-
clickbait
1573      Strong earthquake jolts Islamabad  not-
clickbait
1574      not-
clickbait

[1575 rows x 2 columns]

```

Combining the datasets

```

combined_dataset = pd.concat([positive_dataset, negative_dataset],
ignore_index=True)
df=pd.DataFrame(combined_dataset)

df.to_csv("combined.csv",index=False) #

#reading the first 8 rows of the combined dataset
df.head(8)

```

	baits	Label
0	Man repairs fence to contain dog, hilarity ens...	clickbait
1	Long-Term Marijuana Use Has One Crazy Side Eff...	clickbait
2	The water from his ear trickles into the bucke...	clickbait
3	You'll Never Guess What Nick Jonas Does in the...	clickbait
4	How Cruise Liners Fill All Their Unsold Cruise...	clickbait
5	Could Queen Elizabeth Veto Brexit?	clickbait
6	This Is the Worst Color to Paint Your Kitchen	clickbait
7	The Shocking Truth About Sugar	clickbait

Shuffling the combined dataset using numpy

```

# Create an array of indices from 0 to the length of the combined
dataset
shuffled_indices = np.arange(len(combined_dataset))
# Shuffle the array of indices randomly using numpy's random.shuffle

```

```

function
np.random.shuffle(shuffled_indices)
# Use the shuffled indices to rearrange the rows of the combined
dataset, creating the shuffled dataset
shuffled_dataset = combined_dataset.iloc[shuffled_indices]
df2=pd.DataFrame(shuffled_dataset)
df2.to_csv("shuffled.csv",index=False)

```

```

#reading the first ten rows of the shuffled_dataset
df=pd.DataFrame(shuffled_dataset)
df.head(10)

```

	baits	Label
356	You Won't Believe What Material These Sunglass...	clickbait
1671	Lenovo caught installing adware on new computers	not-clickbait
60	How to tell if someone is a narcissist with on...	clickbait
1339	Disney Donates \$1 Million to Orlando Shooting ...	not-clickbait
1568	Who's minding the marijuana? Banned pesticide ...	not-clickbait
570	Why did Croatia fans disrupt their Euro 2016 m...	clickbait
1214	Don't laugh: Google's Parsey McParseface is a...	not-clickbait
309	Paul McCartney Reveals What Really Split Up th...	clickbait
2113	Gay Chinese man sues mental hospital for tryin...	not-clickbait
538	Abercrombie & Fitch is facing a terrifying rea...	clickbait

Split the shuffled dataset into train, validation, and test sets

```

train_data_percentage = 0.72 # 72% for training
validation_data_percentage = 0.08 # 8% for validation
test_data_percentage = 0.20 # 20% for testing

# Calculate the number of samples for each split
total_samples = len(shuffled_dataset) #number of rows in the combined
dataset
train_data_samples = int(train_data_percentage * total_samples)
#number of rows in the training dataset
validation_data_samples = int(validation_data_percentage *
total_samples) #number of rows in the validation dataset

```

```
test_data_samples = total_samples - train_data_samples -  
validation_data_samples #number of rows in the testing dataset
```

```
#printing the number of rows/ samples in each dataset  
print("***** OUTPUT *****\n")  
print(f"total sample or the number of rows ={total_samples}  
samples/rows")  
print(f"validation samples or the number of rows  
={validation_data_samples} samples/rows")  
print(f"Training samples or the number of rows ={train_data_samples}  
samples/rows")  
print(f"test samples or the number of rows ={test_data_samples}  
samples/rows")
```

```
***** OUTPUT *****
```

```
total sample or the number of rows =2390 samples/rows  
validation samples or the number of rows =191 samples/rows  
Training samples or the number of rows =1720 samples/rows  
test samples or the number of rows =479 samples/rows
```

```
# Splitting the dataset into training and the remaining data  
(remaining_data)  
training_data, remaining_data = train_test_split(shuffled_dataset,  
test_size=(1 - train_data_percentage))
```

```
# Split the remaining data into validation and test sets  
validation_data, test_data = train_test_split(remaining_data,  
test_size=test_data_percentage / (test_data_percentage +  
validation_data_percentage))
```

```
#saving the datasets
```

```
train_dataset=pd.DataFrame(training_data) #training dataset  
train_dataset.to_csv("training_data.csv",index=False)
```

```
validating_dataset=pd.DataFrame(validation_data) #validating set  
validating_dataset.to_csv("validating_data.csv",index=False)
```

```
testing_dataset=pd.DataFrame(test_data) #testing set  
testing_dataset.to_csv('testing_dataset.csv',index=False)
```

```
#reading from each and every dataset after split
```

train_dataset *#reading from the training dataset*

	baits	Label
1187	It is rare for a new animal species to emerge ...	not-clickbait
1701	Ontario parents who object to vaccines could b...	not-clickbait
1673	Panama Papers: British Banker Funded North Kor...	not-clickbait
1781	Burnt: 220 hives containing 250,000 bees	not-clickbait
2014	Artificial intelligence: 'Homo sapiens will be...	not-clickbait
...
436	The 'Right' Age to Get Married	clickbait
2386	Astronomers Watch a Supernova and See Reruns	not-clickbait
2324	German spy agency says ISIS sending fighters d...	not-clickbait
992	The disappeared: Chicago police detain America...	not-clickbait
1702	Pregnant women warned not to travel to Rio Oly...	not-clickbait

[1720 rows x 2 columns]

#reading from the testing dataset

testing_dataset

	baits	Label
1873	Only 3 northern white rhinos left on Earth	not-clickbait
701	Trump's campaign cycles \$6 million into Trump ...	clickbait
2366	Top UN Official Says 'Global War on Terror' Is...	not-clickbait
1698	Kerry calls for democracy in Cuba as U.S. flag...	not-clickbait
1032	Knife turned into police allegedly found on O....	not-clickbait
...
1456	Proposed California Ballot Initiative That Wou...	not-clickbait
22	Here's what those floaty things in your eyes are	clickbait
2168	FIFA scandal: Sepp Blatter wins another term a...	not-clickbait
655	A Waterfall in the Middle of a Lake	clickbait
2131	Ikea vows to be net exporter of renewable ener...	not-clickbait

[479 rows x 2 columns]

#reading from the validation dataset

validating_dataset

	baits	Label
382	Why Has Mark Zuckerberg Put Tape Over His Camera?	clickbait
628	Dad And Son Are Seconds From Assassination By ...	clickbait
1125	Martin Shkreli fired as CEO of KaloBios Pharma...	not-clickbait
1935	Former Brazilian soccer star: Don't come to th...	not-clickbait
363	Guess how much Google paid the guy who briefly...	clickbait

```

...
2248 Stop refugees or we'll stop aid, Germany tells... not-clickbait
574 Guess Who's Complaining About Obama's New Over... clickbait
2042 Saudi Arabia insists UN keeps LGBT rights out ... not-clickbait
1321 Napolitano Says 'We Don't Have To Listen To Th... not-clickbait
1123 Vladimir Putin says the Panama Papers are part... not-clickbait

[191 rows x 2 columns]

```

Calculating the "target rate" for each dataset (training, validation and test)

```

train_data_target_rate = (train_dataset['Label'] ==
'clickbait').mean()
validation_data_target_rate = (validating_dataset['Label'] ==
'clickbait').mean()
test_data_target_rate = (testing_dataset['Label'] ==
'clickbait').mean()

```

what % of the three datasets is it labeled as clickbait?

```

print("***** OUTPUT *****\n")
print(f"{train_data_target_rate.round(4)*100}% of the training data is
labeled as clickbait")
print(f"{validation_data_target_rate.round(4)*100}% of the validating
data is labeled as clickbait")
print(f"{test_data_target_rate.round(4)*100}% of the testing data is
labeled as clickbait")

```

```

***** OUTPUT *****

```

```

34.36% of the training data is labeled as clickbait
35.08% of the validating data is labeled as clickbait
32.78% of the testing data is labeled as clickbait

```

***** PROBLEM 3 – Training a single Bag-of-Words (BOW) Text Classifier

```
# Loading the training and validation datasets
validating_dataset #the dataset containing the validating data
train_dataset #the dataset containing the training data

# Map labels to binary values (1 for clickbait, 0 for non-clickbait)
train_dataset['Label'] = (train_dataset['Label'] ==
'clickbait').astype(int)
validating_dataset['Label'] = (validating_dataset['Label'] ==
'clickbait').astype(int)

# Creating a Pipeline with CountVectorizer and MultinomialNB
pipeline = Pipeline([
    ('vectorizer', CountVectorizer(ngram_range=(1, 2))),
    ('classifier', MultinomialNB())
])

# Fitting the classifier on the training dataset
pipeline.fit(train_dataset['baits'], train_dataset['Label'])

# Predict on the training and validation datasets
train_predictions = pipeline.predict(train_dataset['baits'])
validation_predictions = pipeline.predict(validating_dataset['baits'])

# Computing precision, recall, and F1-score for training and
validation datasets with zero_division='warn'
train_precision = precision_score(train_dataset['Label'],
train_predictions, pos_label=1, zero_division='warn')
train_recall = recall_score(train_dataset['Label'], train_predictions,
pos_label=1, zero_division='warn')
train_f1 = f1_score(train_dataset['Label'], train_predictions,
pos_label=1, zero_division='warn')

validation_precision = precision_score(validating_dataset['Label'],
validation_predictions, pos_label=1, zero_division='warn')
validation_recall = recall_score(validating_dataset['Label'],
```



```
validation_predictions, pos_label=1, zero_division='warn')
validation_f1 = f1_score(validating_dataset['Label'],
validation_predictions, pos_label=1, zero_division='warn')
```

```
# Print the results
```

```
print("***** OUTPUT *****\n")
print(f"Training Precision is {train_precision} or
{train_precision.round(4)*100}%")
print(f"Training Recall is {train_recall} or
{train_recall.round(4)*100}%")
print(f"Training F1-score is {train_f1} or
{train_recall.round(4)*100}%")
print("\n")
print(f"Validation Precision is {validation_precision} or
{validation_precision.round(4)*100}% ")
print(f"Validation Recall is {validation_recall} or
{validation_precision.round(4)*100}%")
print(f"Validation F1-score is { validation_f1} or
{validation_f1.round(4)*100}%")
```

```
***** OUTPUT *****
```

```
Training Precision is 0.9915966386554622 or 99.16%
Training Recall is 0.9983079526226735 or 99.83%
Training F1-score is 0.9949409780775718 or 99.83%
```

```
Validation Precision is 0.8615384615384616 or 86.15%
Validation Recall is 0.835820895522388 or 86.15%
Validation F1-score is 0.8484848484848485 or 84.85000000000001%
```

***** PROBLEM 4 –

Hyperparameter Tuning*****

```
# Mapping labels to binary values (1 for clickbait, 0 for non-
clickbait)
```

```

train_dataset['Label'] = (train_dataset['Label'] ==
'clickbait').astype(int)
validating_dataset['Label'] = (validating_dataset['Label'] ==
'clickbait').astype(int)

# Defining a grid of hyperparameters to search
parameter_grid = {
    'vectorizer__max_df': [0.5, 0.75, 1.0], # Vary max_df for
CountVectorizer
    'classifier__alpha': [1.0, 0.5, 0.1], # Vary smoothing for
MultinomialNB
    'vectorizer__ngram_range': [(1, 1), (1, 2)] # Include or exclude
bigrams in CountVectorizer
}

# Initializing an empty list to store the results
results = []

# Iterating over the parameter grid
for params in ParameterGrid(parameter_grid):
    # Creating a Pipeline with CountVectorizer and MultinomialNB with
the current parameters
    pipeline = Pipeline([
        ('vectorizer',
CountVectorizer(max_df=params['vectorizer__max_df'],
ngram_range=params['vectorizer__ngram_range'])),
        ('classifier',
MultinomialNB(alpha=params['classifier__alpha']))
    ])

    # Fitting the classifier on the training dataset
    pipeline.fit(train_dataset['baits'], train_dataset['Label'])

    # Predicting on the validation dataset
    validation_predictions =
pipeline.predict(validating_dataset['baits'])

    # Computing precision, recall, and F1-score for validation
dataset
    validation_precision =
precision_score(validating_dataset['Label'],
validation_predictions, zero_division=1)
    validation_recall = recall_score(validating_dataset['Label'],
validation_predictions, zero_division=1)

```

```

validation_f1 = f1_score(validating_dataset['Label'],
validation_predictions,zero_division=1)

# Storing the results
results.append({
    'params': params,
    'validation_precision': validation_precision,
    'validation_recall': validation_recall,
    'validation_f1': validation_f1
})

# Converting results to a DataFrame for analysis
results_dataframe = pd.DataFrame(results)

# Sorting the results by F1-score in descending order
results_dataframe = results_dataframe.sort_values(by='validation_f1',
ascending=False)

# Displaying the top and bottom results
print("***** OUTPUT *****\n")
print("Top Results")
print(results_dataframe.head())
print("Bottom Results")
print(results_dataframe.tail())

```

***** OUTPUT *****

Top Results

	params
validation_precision \	
0	{'classifier__alpha': 1.0, 'vectorizer__max_df...
1.0	
1	{'classifier__alpha': 1.0, 'vectorizer__max_df...
1.0	
16	{'classifier__alpha': 0.1, 'vectorizer__max_df...
1.0	
15	{'classifier__alpha': 0.1, 'vectorizer__max_df...
1.0	
14	{'classifier__alpha': 0.1, 'vectorizer__max_df...
1.0	

	validation_recall	validation_f1
0	1.0	1.0
1	1.0	1.0
16	1.0	1.0

15	1.0	1.0
14	1.0	1.0

Bottom Results

	params
validation_precision \	
5	{'classifier__alpha': 1.0, 'vectorizer__max_df...
1.0	
4	{'classifier__alpha': 1.0, 'vectorizer__max_df...
1.0	
3	{'classifier__alpha': 1.0, 'vectorizer__max_df...
1.0	
2	{'classifier__alpha': 1.0, 'vectorizer__max_df...
1.0	
17	{'classifier__alpha': 0.1, 'vectorizer__max_df...
1.0	

	validation_recall	validation_f1
5	1.0	1.0
4	1.0	1.0
3	1.0	1.0
2	1.0	1.0
17	1.0	1.0

*****PROBLEM 5

– Model selection *****

```
#To select the best model from the results of the hyperparameter tuning,
#i choose the one that achieved the highest F1-score on the validation set. how do choose.

# Finding the best model based on validation F1-score values
best_model = max(results, key=lambda x: x['validation_f1'])

# Displaying the best model's parameters and validation F1-score value
print("***** OUTPUT *****\n")
print("Best Model - Parameters ")
print(best_model['params'])
print("\n")
print(" Validation F1-Score value of the Best Model is ",
best_model['validation_f1'])
```

***** OUTPUT *****

Best Model - Parameters

```
{'classifier__alpha': 1.0, 'vectorizer__max_df': 0.5,  
'vectorizer__ngram_range': (1, 1)}
```

Validation F1-Score value of the Best Model is 1.0

*#applying the model to my test set and computing the precision,
recall, and F1-score values*

```
# Creating a pipeline with the best model's parameters  
best_model_pipeline = Pipeline([  
    ('vectorizer', CountVectorizer(max_df=best_model['params']  
    ['vectorizer__max_df'], ngram_range=best_model['params']  
    ['vectorizer__ngram_range'])),  
    ('classifier', MultinomialNB(alpha=best_model['params']  
    ['classifier__alpha']))  
])
```

```
# Fiting my best model on the training dataset  
best_model_pipeline.fit(train_dataset['baits'],  
train_dataset['Label'])
```

```
# Predicting on the test dataset  
test_predictions =  
best_model_pipeline.predict(testing_dataset['baits'])
```

```
#removing posible NAN values  
clean_testing_data = testing_dataset.dropna()
```

```
# Convert true labels to numeric format  
clean_testing_data['Label'] =  
clean_testing_data['Label'].map({'clickbait': 1, 'not-clickbait': 0})
```

```
# Computing the precision, recall, and F1-score values for the test  
dataset  
test_precision = precision_score(clean_testing_data['Label'],  
test_predictions, zero_division=1)  
test_recall = recall_score(clean_testing_data['Label'],  
test_predictions)  
test_f1 = f1_score(clean_testing_data['Label'], test_predictions)
```

```
print("***** OUTPUT *****\n")
```

```
# Display the test set results
print("These are Test Set Metrics for Best Model ")
print("Precision ", test_precision)
print("Recall ", test_recall)
print("F1-Score " , test_f1)
```

***** OUTPUT *****

```
These are Test Set Metrics for Best Model
Precision  1.0
Recall    0.0
F1-Score  0.0
```

*****PROBLEM 6 – Key Indicators*****

```
#Convert the text to lowercase to ensure consistency.
train_dataset['bait'] = train_dataset['bait'].str.lower()

# Defining a dictionary to map labels to binary values
label_mapping = {'clickbait': 1, 'not-clickbait': 0} # Mapping labels
to binary values (1 for clickbait, 0 for non-clickbait)

# Using the map method to update my 'Label' column
train_dataset['Label'] = train_dataset['Label'].map(label_mapping)

# Creating a CountVectorizer to extract key words
vectorizer = CountVectorizer(ngram_range=(1, 2))
```

```

# Transforming the training data into a feature matrix
X_train = vectorizer.fit_transform(train_dataset['bait'])

# Initializing and training a Multinomial Naive Bayes classifier
clf = MultinomialNB()
clf.fit(X_train, train_dataset['Label'])

# Getting the feature log probabilities for the "clickbait" class
clickbait_log_probs = clf.feature_log_prob[1]

# Getting the corresponding vocabulary words
feature_names = np.array(vectorizer.get_feature_names_out())

# Creating a DataFrame to analyze the results
clickbait_dataframe = pd.DataFrame({'Word': feature_names, 'Log
Probability': clickbait_log_probs})

# Sorting the DataFrame by log probability in descending order
clickbait_dataframe = clickbait_dataframe.sort_values(by='Log
Probability', ascending=False)

# Display the top 5 words with the highest log probability
top_clickbait_words = clickbait_dataframe.head(6)
print("***** OUTPUT *****\n")
print("Top 5 Clickbait Indicator Words")
print(top_clickbait_words)

```

***** OUTPUT *****

Top 5 Clickbait Indicator Words

	Word	Log Probability
16995	the	-4.851182
20060	you	-5.373704
17625	to	-5.462499
17409	this	-5.521340

9028	is	-5.567860
11905	of	-5.789021

***** PROBLEM 7 – Regular expression

```
# A list of words that are recognized as clickbaits in the previous
Question
top_keywords = top_clickbait_words['Word'].unique()

# Constructing a regular expression pattern to match any of the top
keywords with word boundaries
pattern = r'\b(?:' + '|'.join(re.escape(keyword) for keyword in
top_keywords) + r')\b'

#A testing message
text = "you will be a billionaire within two days if you do the
following. Follow this steps ."
print("***** OUTPUT *****\n")
# Using re.search function to find matches in the text
if re.search(pattern, text):
    print("Keyword found in the text.")
else:
    print("Keyword not found in the text.")
```

```
***** OUTPUT *****
```

```
Keyword found in the text.
```



```

#removing any nans
testing_dataset=testing_dataset.dropna()

# Defining a dictionary to map labels to binary values
label_mapping = {'clickbait': 1, 'not-clickbait': 0} # Mapping labels
to binary values (1 for clickbait, 0 for non-clickbait)

# Using the map method to update my 'Label' column
testing_dataset['Label'] = testing_dataset['Label'].map(label_mapping)

# Creating a function to check if any top keywords are present in the
text from the test data_set
def keyword_classifier(text):
    pattern = r'\b(?:' + '|'.join(re.escape(keyword) for keyword in
top_keywords) + r')\b'
    return bool(re.search(pattern, text))

# Applying the function to test_dataset
testing_dataset['Predicted'] =
testing_dataset['baits'].apply(keyword_classifier)

# Calculating the precision and recall values

precision = precision_score(testing_dataset['Label'],
testing_dataset['Predicted'], zero_division='warn')

recall = recall_score(testing_dataset['Label'],
testing_dataset['Predicted'], zero_division='warn')

f1=f1_score(testing_dataset['Label'], testing_dataset['Predicted'],
zero_division='warn')

```

```
#printing the results
print("***** OUTPUT *****\n")
print(f"The Precision of this classifier is {precision} or  
{precision.round(4)*100}%")
print(f"The Recall of this classifier is {recall} or  
{recall.round(4)*100}%")
print(f"The Recall of this classifier is {f1} or  
{f1.round(4)*100}%")

***** OUTPUT *****

The Precision of this classifier is 0.32589285714285715 or 32.59%
The Recall of this classifier is 0.46496815286624205 or 46.5%
The Recall of this classifier is 0.3832020997375329 or 38.32%
```